


Local Features

261458 & 261753 Computer Vision

#9

Global vs. Local Features

Feature extract from entire image = **Global Features**




Consider the global features extracted from these images

261458 & 261753 Computer Vision

#9

Global vs. Local Features

Feature extract from some parts of image = **Local Features**



Consider the local features extracted from these images

261458 & 261753 Computer Vision

#9

Advantages Local Features

- Locality:
 - Features are local, so robust to occlusion and clutter
- Distinctiveness:
 - Can differentiate a large database of objects
- Quantity
 - Hundreds or thousands in a single image
- Efficiency
 - Real-time performance achievable
- Generality
 - Exploit different types of features in different situations

261458 & 261753 Computer Vision

#9


Applications Using Local Features

5

- Image alignment (e.g., image stitching)
- Object recognition
- Motion tracking
- 3D reconstruction
- Depth estimation stereo images
- Indexing and database retrieval
- Robot navigation
- ... other


Image Stitching

6



Object Recognition

7



By David Lowe


Image Stitching

8



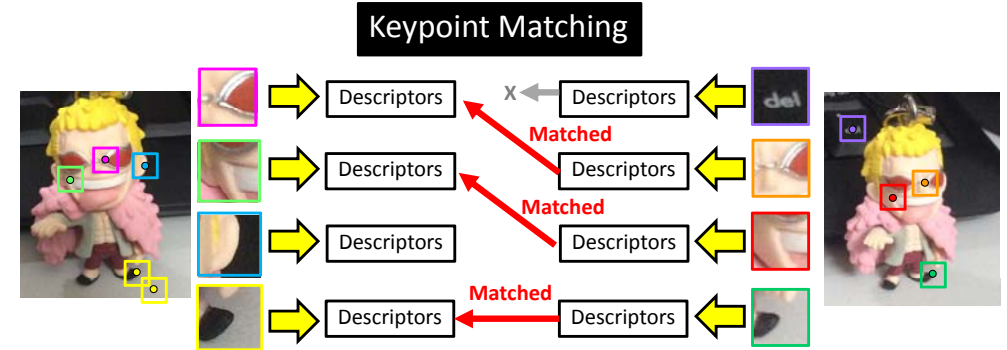
<http://www.cs.toronto.edu/~kyros/courses/2503/Handouts/features.pdf>

Local Feature Based Methods



- Feature Point Detection (Keypoint Detection)
- Feature Descriptor Extraction
- Feature Point Matching


Local Feature Based Methods



- Feature Point Detection (Keypoint Detection)
- Feature Descriptor Extraction
- Feature Point Matching


Local Features

- Find features that are invariant to transformations
 - Geometric invariance: translation, rotation, scale
 - Photometric invariance: brightness, exposure,...



Feature Points [Keypoints]

- Suppose we only consider a small window of pixels
 - What define whether a feature point is a good or bad candidate?



Slide adapted from Darya Frolova, Denis Simakov, Weizmann Institute.

Feature Points [Keypoints]

13



Local Measures of Uniqueness


14

Local measure of "feature uniqueness"

- How does the window change when you shift it?
- Shifting the window in *any* direction causes a big change


"Flat" region:

no change in all directions




"Edge":

no change along the edge direction



"Corner":

significant change in all directions




Slide adapted from Darya Frolova, Denis Simakov, Weizmann Institute.

Keypoint Detection

15

Consider shifting the window W by (u, v)

- How do the pixels in W change?
- Compare each pixel before and after by summing up the squared differences (SSD)
- This defines an SSD "error" of $E(u, v)$



$$E(u, v) = \sum_{(x, y) \in W} [I(x + u, y + v) - I(x, y)]^2$$

Slide adapted from Darya Frolova, Denis Simakov, Weizmann Institute.

Keypoint Detection

16

Taylor series expansion of I :

$$I(x+u, y+v) = I(x, y) + \frac{\partial I}{\partial x}u + \frac{\partial I}{\partial y}v + \text{higher order terms}$$

- If the motion (u, v) is small, then first order approximation is good

$$\begin{aligned} I(x+u, y+v) &\approx I(x, y) + \frac{\partial I}{\partial x}u + \frac{\partial I}{\partial y}v \\ &\approx I(x, y) + [I_x \ I_y] \begin{bmatrix} u \\ v \end{bmatrix} \end{aligned}$$

shorthand: $I_x = \frac{\partial I}{\partial x}$

- Plugging this into the formula on the previous slide...

Keypoint Detection

17

$$\begin{aligned}
 E(u, v) &= \sum_{(x, y) \in W} [I(x + u, y + v) - I(x, y)]^2 \\
 &\approx \sum_{(x, y) \in W} [I(x, y) + [I_x \ I_y] \begin{bmatrix} u \\ v \end{bmatrix} - I(x, y)]^2 \\
 &\approx \sum_{(x, y) \in W} \left[[I_x \ I_y] \begin{bmatrix} u \\ v \end{bmatrix} \right]^2 \\
 &\approx \sum_{(x, y) \in W} [u \ v] \begin{bmatrix} I_x^2 & I_x I_y \\ I_y I_x & I_y^2 \end{bmatrix} \begin{bmatrix} u \\ v \end{bmatrix} \\
 E(u, v) &\approx [u \ v] \begin{bmatrix} \sum_{(x, y) \in W} I_x^2 & \sum_{(x, y) \in W} I_x I_y \\ \sum_{(x, y) \in W} I_x I_y & \sum_{(x, y) \in W} I_y^2 \end{bmatrix} \begin{bmatrix} u \\ v \end{bmatrix}
 \end{aligned}$$


Keypoint Detection

18

$$H = \begin{bmatrix} \sum_{(x, y) \in W} I_x^2 & \sum_{(x, y) \in W} I_x I_y \\ \sum_{(x, y) \in W} I_x I_y & \sum_{(x, y) \in W} I_y^2 \end{bmatrix} \Rightarrow \begin{bmatrix} \mathbf{x}_+, \lambda_+ \\ \mathbf{x}_-, \lambda_- \end{bmatrix} \quad (\lambda_+ \geq \lambda_-)$$


Eigenvalues and eigenvectors of H

- \mathbf{x}_+ = Direction $\begin{bmatrix} u \\ v \end{bmatrix}$ of largest increase in $E(u, v)$
- \mathbf{x}_- = Direction $\begin{bmatrix} u \\ v \end{bmatrix}$ of smallest increase in $E(u, v)$
- λ_+ = Amount of increasing in direction x_+
- λ_- = Amount of increasing in direction x_-



Keypoint Detection

19




Keypoint Detection

20

Here's what to do

- Compute the gradient at each point in the image
- Create the H matrix from the entries in the gradient
- Compute the eigenvalues
- Find **points with large response** ($\lambda_- > \text{threshold}$)
- Choose those points where λ_- is a local maximum as **feature points**



Other Corner Measure Functions

21

Harris Detector

$$M_H = \det(H) - \alpha [\text{trace}(H)]^2$$

$$= (1 + 2\alpha)\lambda_- \lambda_+ + \alpha(\lambda_-^2 + \lambda_+^2)$$

α is a tuning parameter(= 0.04 in the original paper)

Noble Detector


$$M_N = \frac{\det(H)}{\text{trace}(H) + \varepsilon} = \frac{\lambda_- \lambda_+}{\lambda_- + \lambda_+ + \varepsilon}$$

ε is a small number to avoid zero division

- $\text{trace}(H) = h_{11} + h_{22}$
- Very similar to λ_- but less expensive (no square root)

Harris Operator

22




Keypoint Detection Example

23




Keypoint Detection Example

24



Keypoint Detection Example

25




261458 & 261753 Computer Vision

#9

Keypoint Detection Example

26



261458 & 261753 Computer Vision

#9

Keypoint Detection Example

27



261458 & 261753 Computer Vision

#9

Keypoint Detection Invariance?

28

- Suppose you **rotate** the image by some angle
 - Will you still get the same features ?
- What if you change the brightness?
- How about scale ?


261458 & 261753 Computer Vision

#9

Scale Invariant Detection

29

- Suppose you're looking for corners



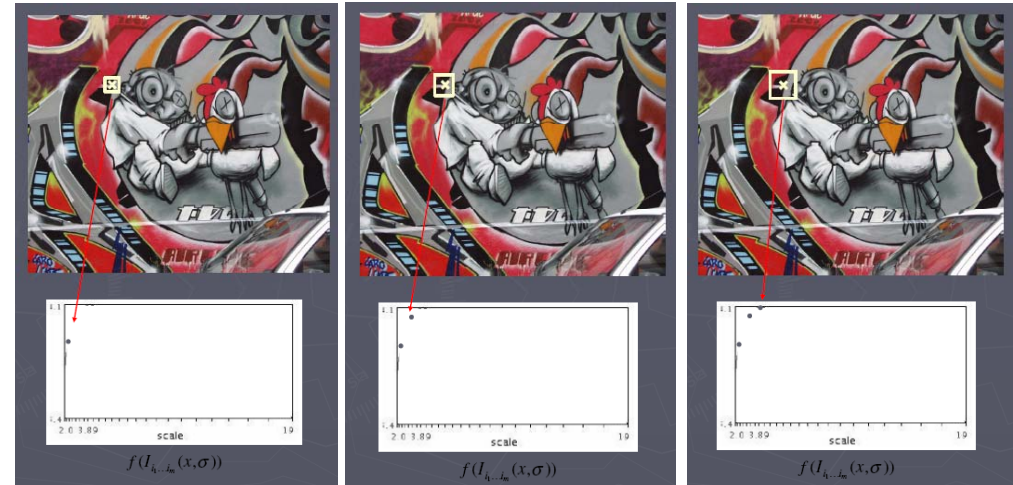
- Key idea: find scale that gives local maximum of f
 - f is a local maximum in both position and scale
 - Common definition of f : Laplacian
(or difference between two Gaussian filtered images with different sigmas)

261458 & 261753 Computer Vision

#9

Automatic Scale Selection

30



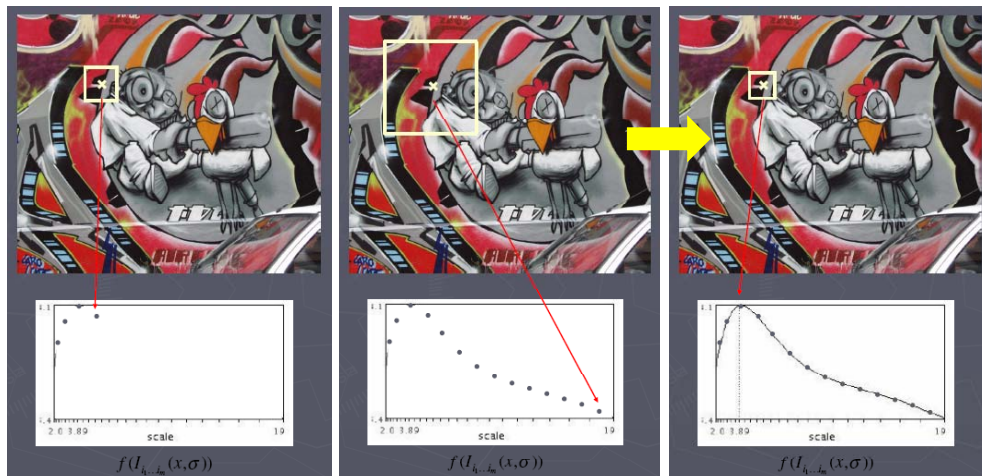
Lindeberg et al., 1996

261458 & 261753 Computer Vision

#9

Automatic Scale Selection

31



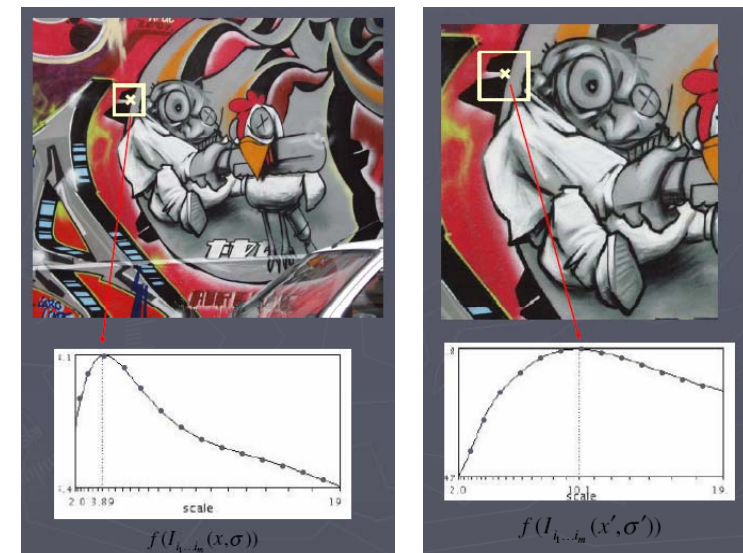
Lindeberg et al., 1996

261458 & 261753 Computer Vision

#9

Automatic Scale Selection

32




261458 & 261753 Computer Vision

#9

Keypoint Descriptors

33

- We know how to detect good points
- Next question: How to match them ?



- Lots of possibilities (this is a popular research area)
 - Simple option: match square windows around the point
 - State of the art approach: SIFT, SURF, BRIEF.

Invariant Descriptors

34

- Suppose we are comparing two images I_1 and I_2
 - I_2 may be a transformed version of I_1
 - What kinds of transformations are we likely to encounter in practice?
- We'd like to find the same feature regardless of the transformation
 - This is called transformational invariance
 - Most feature methods are designed to be invariant to
 - Translation, 2D rotation, scale
 - They can usually also handle
 - Limited 3D rotations (SIFT works up to about 60 degrees)
 - Limited affine transformations (some are fully affine invariant)
 - Limited illumination/contrast changes


Rotation Invariance for Feature Descriptors

35



- Find dominant orientation of the image patch. This is given by \mathbf{x}_+ (the eigenvector of H corresponding to the larger eigenvalue λ_+)
- Rotate the patch according to this angle
- Extract descriptor from the patch

Figure by Matthew Brown



Keypoint Matching

37

- Given a keypoint in I_1 , how to find the best match in I_2 ?
 - Define distance function that compares two descriptors
 - Test all the features in I_2 , find the one with minimum distance

Distance Function

38

- How to define the difference between two features f_1 and f_2 ?
 - Simple approach is $SSD(f_1, f_2)$
 - Sum of square differences between entries of the two descriptors
 - Can give good scores to very ambiguous (bad) matches



Distance Function

39

- How to define the difference between two features f_1 and f_2 ?
 - Better approach: ratio distance = $SSD(f_1, f_2) / SSD(f_1, f_2')$
 - f_2 is the best SSD match to f_1 in I_2
 - f_2' is the 2nd best SSD match to f_1 in I_2
 - Gives small values for ambiguous matches

